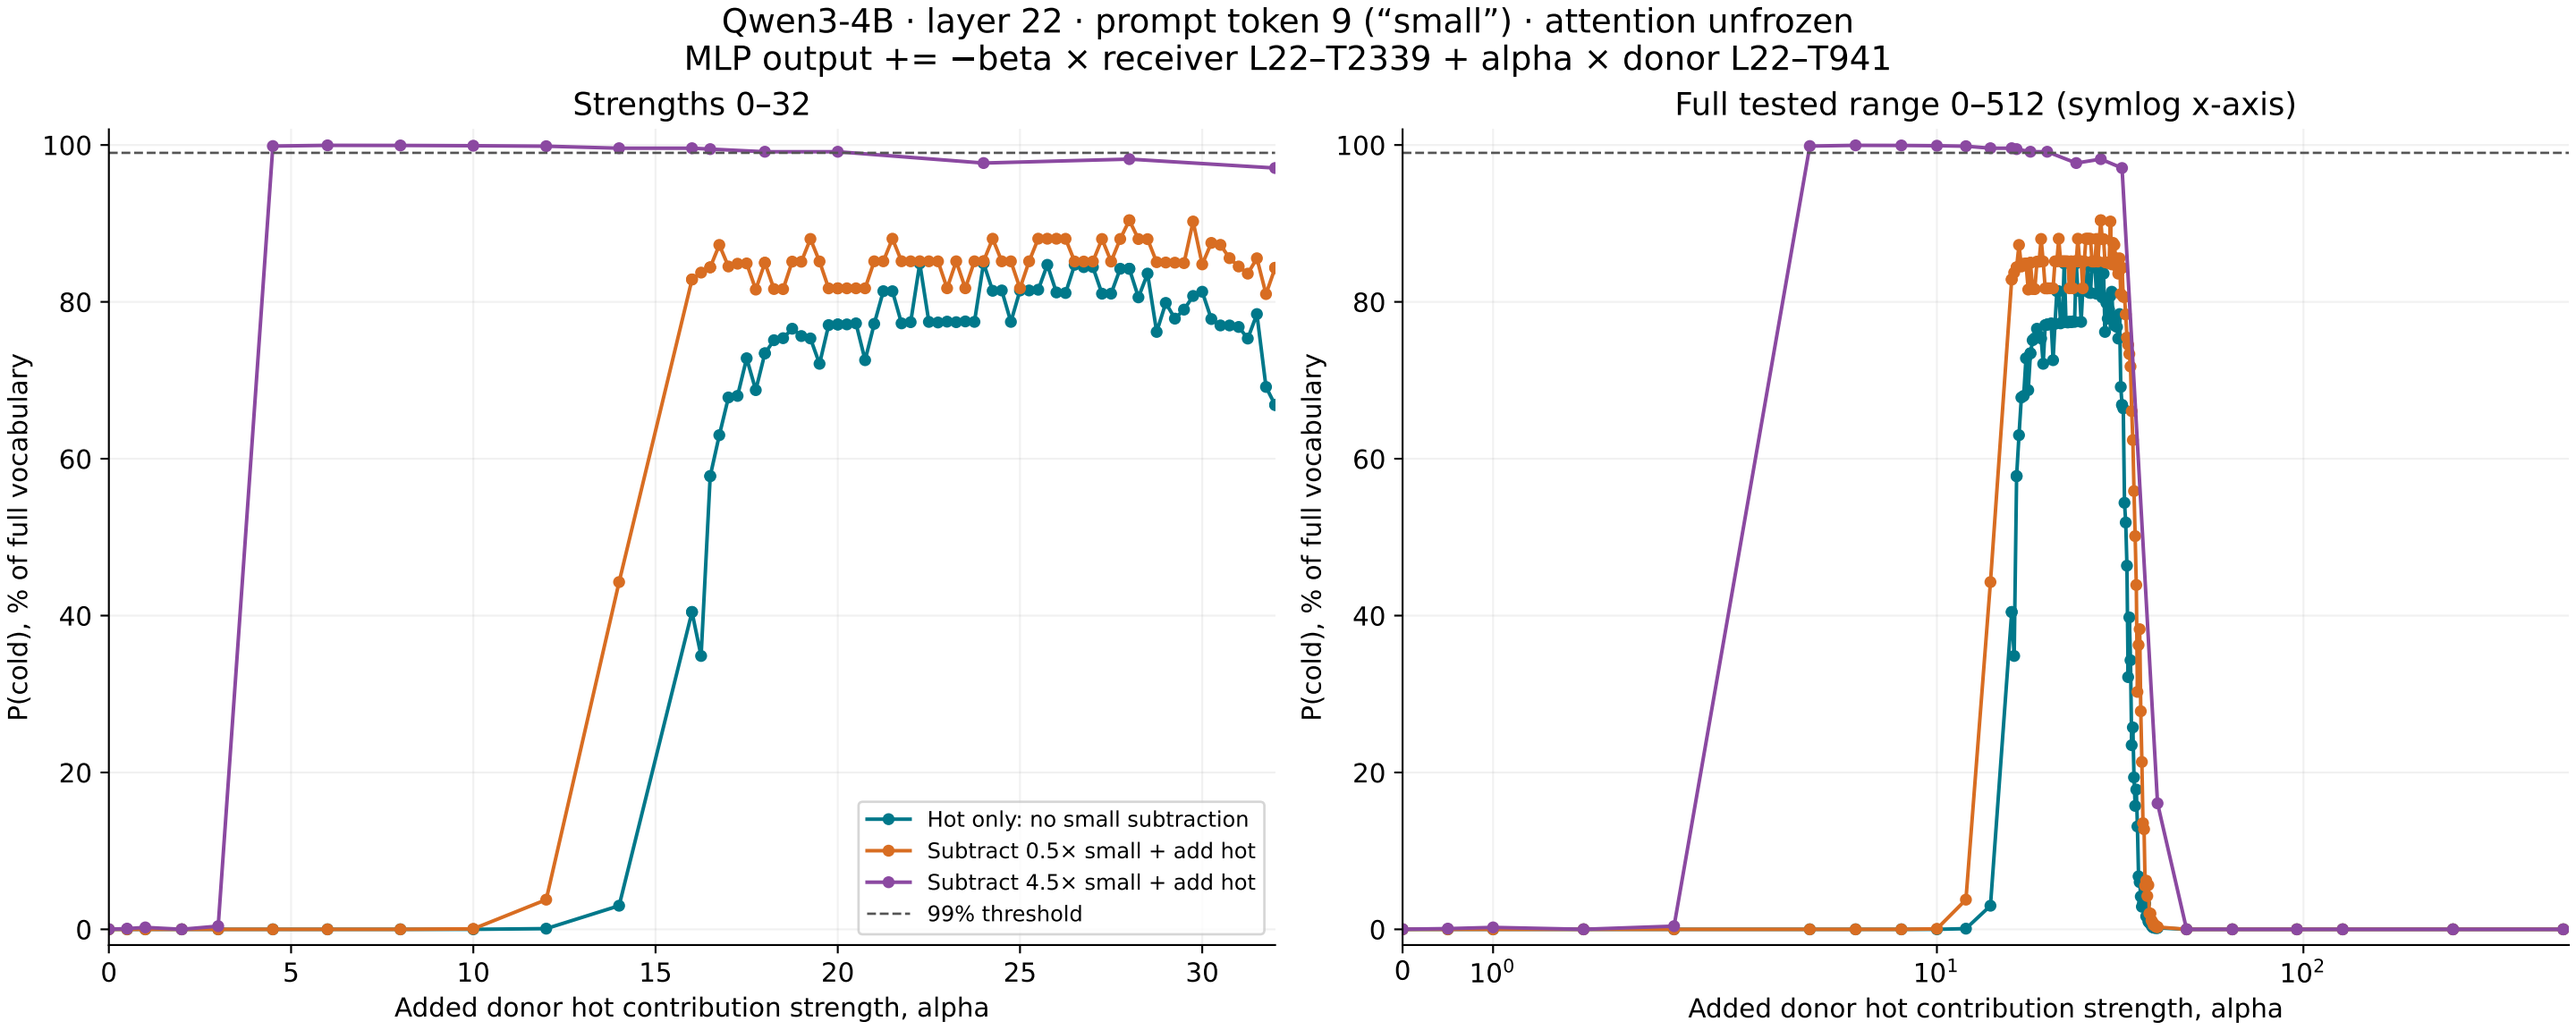
Qwen3-4B · layer 22 · prompt token 9 (“small”) · attention unfrozen
MLP output += −beta × receiver L22–T2339 + alpha × donor L22–T941
Strengths 0–32
Full tested range 0–512 (symlog x-axis)
P(cold), % of full vocabulary
Added donor hot contribution strength, alpha
Hot only: no small subtraction
Subtract 0.5× small + add hot
Subtract 4.5× small + add hot
99% threshold